



Operating Elasticsearch in Kubernetes

microxchg ²⁰¹⁹

MIKKEL LARSEN

[@mikkeloscar](https://twitter.com/mikkeloscar)

OLIVER TROSIEN

[@otrosien](https://twitter.com/otrosien)

2019-04-02



WHO ARE WE?

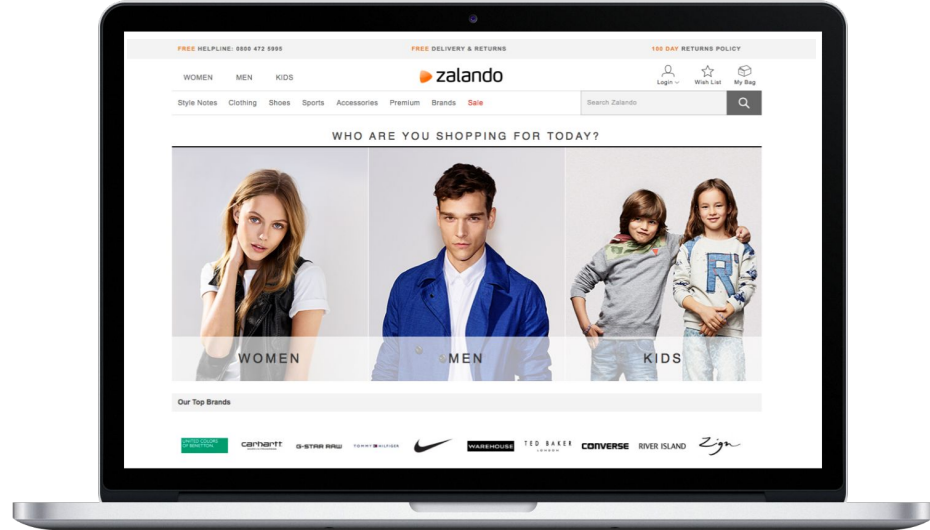



Oliver
Core Search Services
Elasticsearch



Mikkel
Cloud Infrastructure
Kubernetes

“EUROPE’S LEADING ONLINE FASHION PLATFORM”



WE BRING FASHION TO PEOPLE IN 17 COUNTRIES

17 markets


7 fulfillment centers

26 million active customers

5.4 billion € revenue 2018

250 million visits per month

15,000 employees in Europe



SEARCH @ ZALANDO

300k+

Products
per country

45%

Mobile Traffic

~2000

Brands

12K

QPS

~700


Categories

8K


Updates/s


















DATA INGESTION



Kafka / Nakadi *



* nakadi.io




KUBERNETES @ ZALANDO

Default
Deployment
Target

1400~
nodes

Since
Oct 2016




zalando

114
clusters

Node
Autoscaling

From v1.4
to v1.12



Elasticsearch
2.500 vCPUs
1 TB RAM

Elasticsearch in Kubernetes

RUNNING ELASTICSEARCH IN KUBERNETES

1. **Safe automatic updates**
(Including Kubernetes cluster updates)
2. **Advanced auto-scaling for cost efficiency**




UPDATING ELASTICSEARCH (STATEFULSET)

1) PreStop Hook ([bash script](#))

- Exclude node in ES
- Wait for node to drain (up to 1h)
- Data is moved to existing nodes

2) PostStart Hook ([bash script](#))

- Remove all excludes
- Let ES rebalance from existing nodes




ELASTICSEARCH DATA SETS

```
apiVersion: zalando.org/v1
kind: ElasticsearchDataSet
metadata:
  name: test-cluster
spec:
  scaling:
    {...}
  replicas: 3
  template: # PodTemplate
    {...}
  volumeClaimTemplates:
    {...}
```

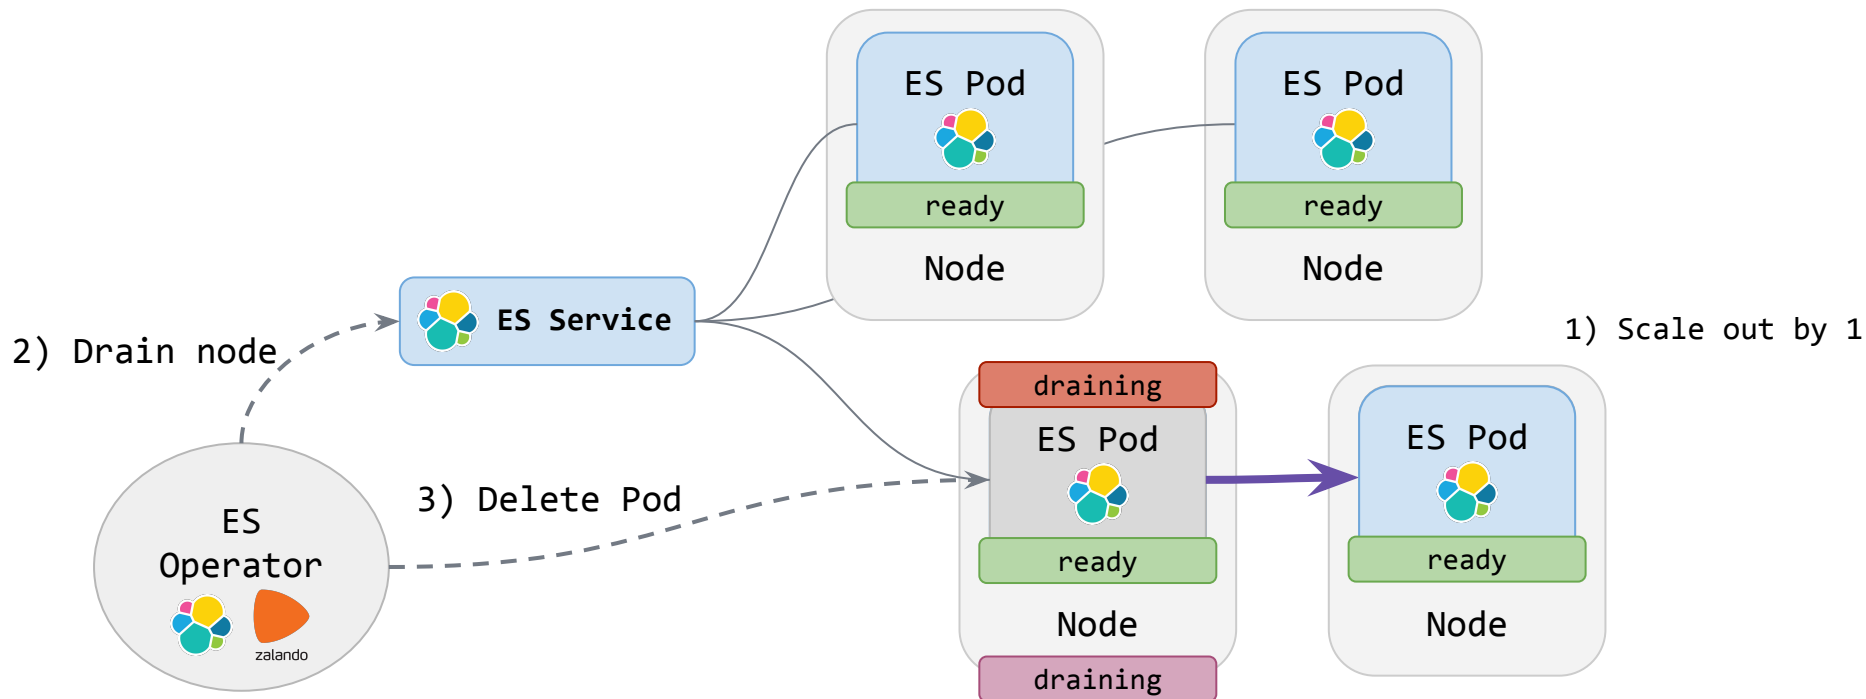
github.com/zalando-incubator/es-operator

ELASTICSEARCH DATA SETS




github.com/zalando-incubator/es-operator

UPDATING ELASTICSEARCH (OPERATOR)



github.com/zalando-incubator/es-operator

SCALING UP ELASTICSEARCH (1)




Thresholds

- CPU
- Duration
- Cooldown

Boundaries


- Max # Pod replicas



SCALING DOWN ELASTICSEARCH

**DON'T OPERATE
WHEN CLUSTER
IS NOT GREEN!**

METRICS







Thresholds

- CPU
- Duration
- Cooldown

Boundaries


- Min # Replica
- Max # Shards per node
- Max disk usage (%)



Decrease Pod replicas →

SCALING UP ELASTICSEARCH (2)

METRICS

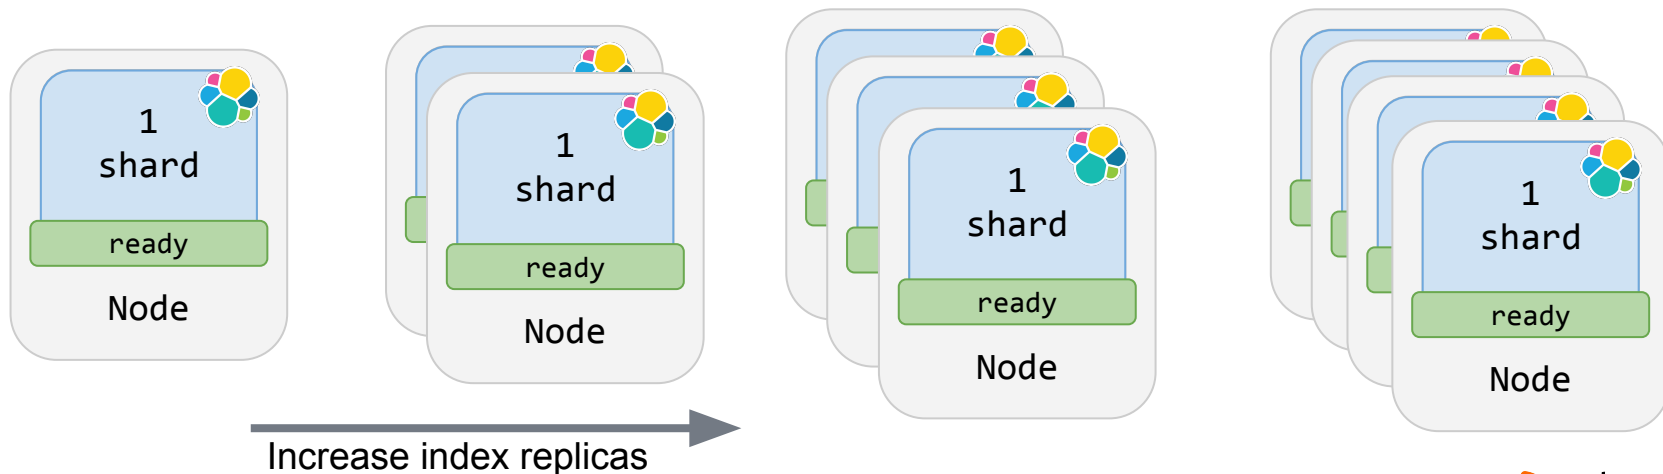


Thresholds

- CPU
- Duration
- Cooldown

Boundaries

- Min # Shards per node
- Max # Pod replicas



DEMO



OPEN SOURCE

ES Operator

github.com/zalando-incubator/es-operator

microXchg demo


github.com/otrosien/microxchg19-demo

Nakadi

nakadi.io

Kubernetes on AWS

github.com/zalando-incubator/kubernetes-on-aws



zalando



MIKKEL LARSEN

mikkel.larsen@zalando.de

[@mikkeloscar](https://www.instagram.com/mikkeloscar)



OLIVER TROSIEN

oliver.trosien@zalando.de

[@otrosien](https://www.instagram.com/otrosien)



2019-04-02