



# POSTGRESQL MONITORING IN ZALANDO

---



## Helsinki PostgreSQL Meetup



October 2019






# POSTGRESQL IN ZALANDO

## POSTGRESQL IN ZALANDO

Zalando works with PostgreSQL since approx. 2010 (running in DC)

- 2015** - migration to AWS (using RDS);  
Patroni and Spilo have been started.
- 2016-2017** - using provided PostgreSQL clusters  
(based on Spilo/STUPS);
- 2018-...** - using PostgreSQL operator.





# POSTGRESQL OPERATOR

When a new postgresql custom resource appears, the operator creates:


1. **StatefulSet** for PostgreSQL/Patroni cluster
2. **Service** for master node (ClusterIP or LB)
3. **Service** for replica nodes (ClusterIP or LB)
4. Extra DNS names for the services *if needed*




kubernetes



Amazon  
**EBS**





## **ZMON - ZALANDO MONITORING SYSTEM**

## Few facts about ZMON

- Created in 2013 during a HackWeek, in production since 2014;
- Optimized for our use-case: autonomous teams, multiple Kubernetes clusters, common central storage;
- Open-sourced (APL2.0) and available from GitHub;
- Stores time-series data in KairosDB (based on top of Cassandra);
- Stores infrastructure data in PostgreSQL.

## What's important for us?

- ZMON uses the “pull” model, its workers *fetch* the metrics
- ZMON is a *distributed* monitoring system, the workers are running in every K8s cluster;
- Autodiscover results in the most up-to-date view of our apps;
- Checks are centralized and may be shared between the teams;
- PostgreSQL Patroni/Spilo clusters are fully supported.



# INFRASTRUCTURE MONITORING



## What are the metrics we need?

- Kubernetes
  - Pods CPU
  - Network I/O
  - Free space in Persistent Volumes
  - Open TCP connections for PostgreSQL processes
- AWS
  - EBS I/O
  - ELB throughput (rarely)
  - Backup S3 bucket



Amazon Cloudwatch

# Prometheus

## Description

Monitor PostgreSQL cluster masters


## Command


```
def check():
    try:
        d = http("http://" + entity["ip"] + ':8080').json()
    except:
        d = http("http://" + entity["ip"] + ':8080').json()

    r = {}
    r["load1"] = d["system_stats"]["load_average"][0]
    r["load5"] = d["system_stats"]["load_average"][1]
    r["connections"] = {
        "total": d["postgresql"]["connections"]["total"],
        "active": d["postgresql"]["connections"]["active"],
    }
    if 'cgroup' in d:
        r["memory"] = d["cgroup"]["memory"]
    else:
        r["memory"] = d["system_stats"]["memory"]["cgroup"]
    r["cpu"] = d["system_stats"]["cpu"]

    return r
```

Interval 1m





## POSTGRESQL INTERNAL METRICS

## How to collect internal metrics

- ZMON worker is running in the same K8s cluster as the PostgreSQL nodes;
- It supports running SQL queries natively;
- We need only to figure out the proper tables/views to select from.



## What are the metrics we need?

- Server
  - idle transactions `pg_stat_activity`
  - failed login attempts
- Tables `pg_stat_user_tables`
  - size / seq\_scans / inserts / updates / deletes
- Indexes `pg_stat_user_indexes`
  - size / scans
- Backups `pg_stat_archiver`
  - WAL archiver status
  - age of last backup *S3 bucket check*

# Idle transactions

## Description

List the pids of the transactions which where idle since 15 or more minutes.

## Command

```
def check():
    query = """
        SELECT *
        FROM   pg_stat_activity
        WHERE  state in ('idle in transaction', 'idle in transaction (aborted)')
        AND   current_timestamp - state_change >= INTERVAL '15' MINUTE
        AND   pid <> pg_backend_pid();
    """

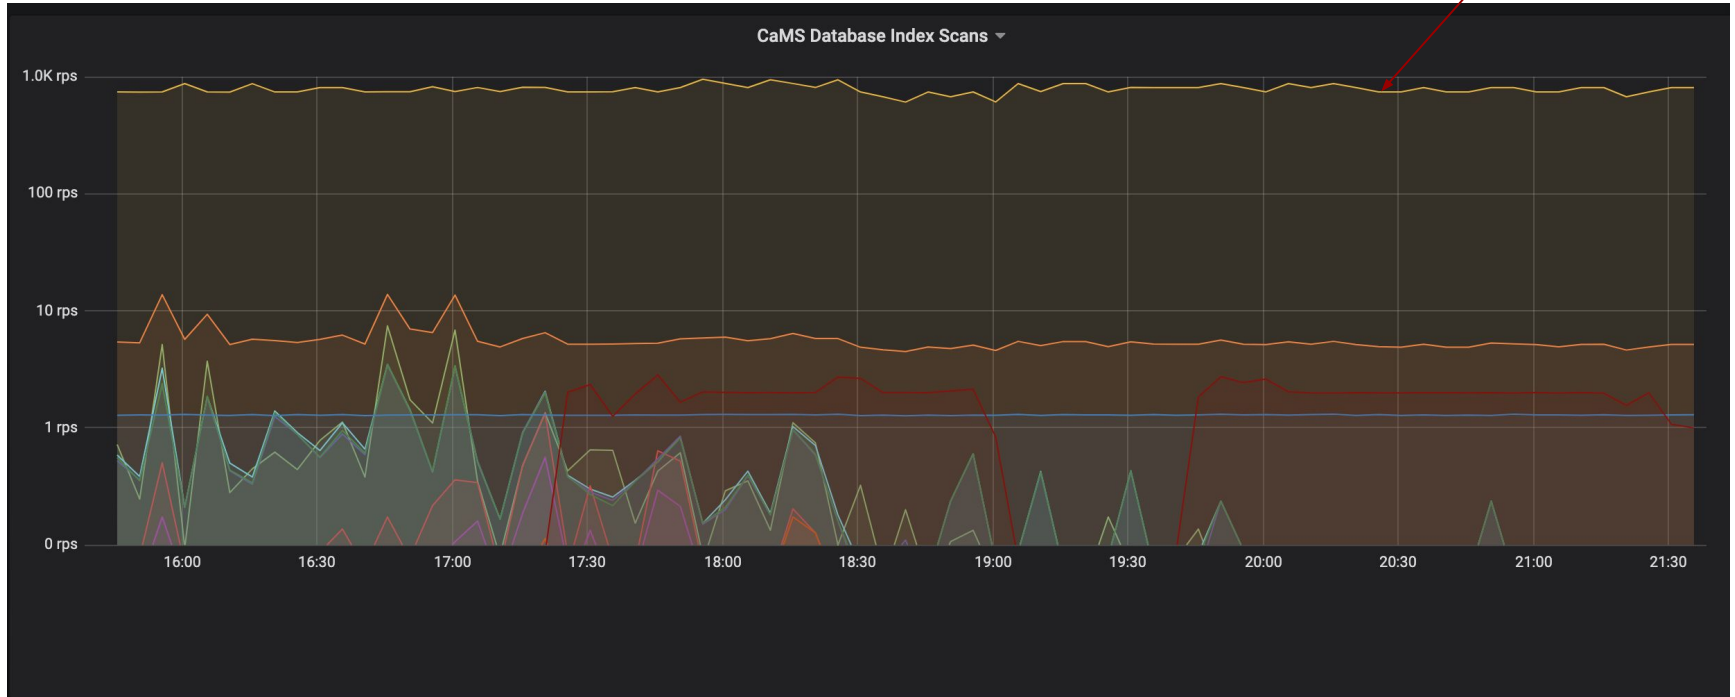
    d1 = {}
    for row in sql(shard='postgres').execute(query).results():
        d1[row['pid']] = {}
        d1[row['pid']]['datname'] = row['datname']
        d1[row['pid']]['username'] = row['username']
        d1[row['pid']]['state_change'] = row['state_change']
        d1[row['pid']]['application_name'] = row['application_name']
        d1[row['pid']]['state'] = row['state']
    return d1
```

Interval 15m

Entities environment = production AND type = postgresql\_cluster

# Index scans

It looks like a problem!



## But what's about authentication?

- ZMON worker uses its own credentials to connect to all the databases in the cluster. The credentials are *separate* from the application credentials;
- ZMON worker user is unique in every K8s cluster: `robot_cluster_name`;
- PostgreSQL role `robot_zmon` is created during the database setup to authorize the ZMON user to access the database objects;
- ZMON may also be authorized to run queries against the application tables or views, but the permissions for the ZMON role should be granted explicitly in DB:

```
GRANT SELECT ON ALL TABLES  
  
IN SCHEMA public TO robot_zmon;
```



# Application metrics

## Description

Retrieves the list of campaignss in error

## Command

```
def check():
    result = {}
    dbResults = sql().execute("""
        SELECT cmp.id, cmp.source_group_id, cmp.destination_id, cmp.runtime_start, cmp.runtime_end,
        cmp.country, cmp.gender, cmp.platform_type, cmp.campaign_type, io.name as io_name, io.destination_id
        as io_destination_id, io.source_id as io_source_id, io.source as io_source
        FROM campaigns.campaign_registry cmp, campaigns.insertion_order_registry io
        WHERE cmp.insertion_order_id=io.id
        AND cmp.is_exported = false
        AND cmp.is_error = true
        AND cmp.runtime_end > now();
    """).results(max_results=10000)

    result['total_count'] = len(dbResults)
    result['items'] = dbResults

    return result
```

Interval 2m

## Key Takeaways

- Get useful metrics from your infrastructure (K8s, AWS, ...)
- Collect PostgreSQL metrics from the `pg_stat_...` views
- Get your application metrics by querying your tables or views directly (but beware of the performance impact)
- Treat your monitoring tool as another application, not as the superuser



**Feel free to reach me:**

**Uri Savelchev**

[<uri.savelchev@zalando.fi>](mailto:uri.savelchev@zalando.fi)

**Find out more about our  
Culture, People & Jobs:**

- **ON SOCIAL MEDIA:**

Linkedin @Zalando SE

Facebook @ Inside Zalando

Instagram @insidezalando

Twitter @ZalandoTech

- **CAREER WEBSITE:** [jobs.zalando.com](https://jobs.zalando.com)

- **CORPORATE WEBSITE:**

[corporate.zalando.com](https://corporate.zalando.com)